

EXAMPLE BREAKAGE

The `acidtest` code (presented in the next section) spawns several child processes that write to a FIFO. Writing happens as a sequence of datums containing the child processes virtual process number `n_vp` and sequence number `k`. The parent process reads from the FIFO and prints out the contents of the datums as they are read. In normal operation of the `acidtest` program, datums are printed by the parent process in a seemingly random order, however, no datum from a single child process is ever out of order. `acidbreak.txt` shows a sample output of the program `acidtest.c` when the write mutex-lock is removed from the FIFO structure (removing lines 14 and 19 from `fifo.c`).

`acidbreak.txt`

```
$ ./acidtest.out
```

...

```
n_vp: 8,      k: 8388
n_vp: 12,     k: 8202
n_vp: 3,      k: 8363
n_vp: 4,      k: 8544
n_vp: 3,      k: 8364
n_vp: 6,      k: 8322
n_vp: 9,      k: 7966
n_vp: 5,      k: 8198
n_vp: 9,      k: 7967
n_vp: 6,      k: 8323
n_vp: 4,      k: 8545
n_vp: 12,     k: 8203
n_vp: 5,      k: 8200
n_vp: 8,      k: 8389
```

...

If we `grep` the output stream of the `acidtest` program for a single virtual process, we can see that the absence of the write mutex-lock causes the datums to be out of order. `acidbreak_grep.txt` shows a sample of the previous `acidtest` output when grepped for virtual process number 13.

`acidbreak_grep.txt`

```
$ ./acidtest.out | grep 'n_vp: 13'
```

...

```
n_vp: 13,     k: 9996
n_vp: 13,     k: 9997
n_vp: 13,     k: 9998
n_vp: 13,     k: 9999
n_vp: 13,     k: 9240
n_vp: 13,     k: 9241
n_vp: 13,     k: 9242
n_vp: 13,     k: 9243
n_vp: 13,     k: 9244
n_vp: 13,     k: 9245
```

...

ACID TEST

acidtest.c

```
1  #include <unistd.h>
2  #include <sys/types.h>
3  #include <sys/mman.h>
4  #include <fcntl.h>
5  #include <errno.h>
6  #include <sys/wait.h>
7  #include <stdio.h>
8  #include <stdlib.h>
9  #include <string.h>
10 #include "fifo.h"
11
12 #define N_ITER 1e4
13 #define N_PROC 16
14 #define MEM_SIZE 0x10000
15
16 int main(int argc, char **argv)
17 {
18     int i,j,k, n_vp;
19     unsigned long d=0;
20     struct fifo *f;
21
22     if( (f=mmap(NULL, MEM_SIZE, PROT_READ|PROT_WRITE, \
23         MAP_SHARED|MAP_ANONYMOUS, -1, 0)) == MAP_FAILED )
24     {
25         perror("mmap error");
26         exit(-1);
27     }
28
29     fifo_init(f);
30
31     for(i=0; i<N_PROC; i++)
32     {
33         switch( fork() )
34         {
35             case -1:
36                 perror("fork error");
37                 exit(-1);
38             case 0:
39                 for(j=0; j<N_ITER; j++)
40                 {
41                     d = (unsigned long)i<<32 | (unsigned long)j;
42                     fifo_wr(f, d);
43                 }
44                 exit(0);
45             default:
46                 break;
47         }
48     }
49     for(j=0; j<N_PROC*N_ITER; j++)
50     {
51         d = fifo_rd(f);
52         n_vp = (int)(d>>32);
53         k = (int)d;
54         printf("n_vp: %d,\t k: %d\n", n_vp, k);
55     }
56     while(wait(NULL)>0);
57     return 0;
58 }
```

FIFO

fifo.h

```
1  #ifndef __FIFO_H
2  #include "sem.h"
3
4  #define MYFIFO_BUFSIZ 0x1000
5  struct fifo
6  {
7      struct sem sfifo_rd;
8      struct sem sfifo_wr;
9      struct sem smutex_rd;
10     struct sem smutex_wr;
11     unsigned long buf[MYFIFO_BUFSIZ];
12     int writr;
13     int rditr;
14 };
15 void fifo_init(struct fifo *f);
16 void fifo_wr(struct fifo *f, unsigned long d);
17 unsigned long fifo_rd(struct fifo *f);
18 #define __FIFO_H
19 #endif
```

fifo.c

```
1  #include "fifo.h"
2
3  void fifo_init(struct fifo *f)
4  {
5      sem_init(&f->sfifo_wr, MYFIFO_BUFSIZ-1);
6      sem_init(&f->sfifo_rd, 0);
7      sem_init(&f->smutex_wr, 1);
8      sem_init(&f->smutex_rd, 1);
9      f->writr=0; f->rditr=0;
10 }
11
12 void fifo_wr(struct fifo *f, unsigned long d)
13 {
14     sem_wait(&f->smutex_wr);
15     sem_wait(&f->sfifo_wr);
16     f->buf[f->writr++] = d;
17     f->writr %= MYFIFO_BUFSIZ;
18     sem_inc(&f->sfifo_rd);
19     sem_inc(&f->smutex_wr);
20 }
21
22 unsigned long fifo_rd(struct fifo *f)
23 {
24     unsigned long ret;
25     sem_wait(&f->smutex_rd);
26     sem_wait(&f->sfifo_rd);
27     ret = f->buf[f->rditr++];
28     f->rditr %= MYFIFO_BUFSIZ;
29     sem_inc(&f->sfifo_wr);
30     sem_inc(&f->smutex_rd);
31     return ret;
32 }
```

SEMAPHORE

sem.h

```
1  #ifndef __SEM_H
2  #include <sys/types.h>
3
4  #define MAX_WAIT 64
5  struct sem
6  {
7      volatile int count;
8      volatile pid_t waitstack[MAX_WAIT];
9      volatile int waitlen;
10     volatile char waitlock;
11     volatile char countlock;
12 };
13 void sem_init(struct sem *s, int count);
14 int sem_try(struct sem *s);
15 void sem_wait(struct sem *s);
16 void sem_inc(struct sem *s);
17 #define __SEM_H
18 #endif
```

sem.c

```
1  #include "tas.h"
2  #include "spinlib.h"
3  #include "sem.h"
4  #include <unistd.h>
5  #include <signal.h>
6  #include <stdio.h>
7  #include <stdlib.h>
8
9  static void handler(int sn){}
10
11 static pid_t pid_pop(struct sem *s)
12 {
13     if(s->waitlen<1)
14         return -1;
15     s->waitlen--;
16     return s->waitstack[s->waitlen];
17 }
18
19 static pid_t pid_push(struct sem *s, pid_t pid)
20 {
21     if(s->waitlen>=MAX_WAIT)
22         return -1;
23     s->waitstack[s->waitlen++] = pid;
24     return 0;
25 }
26
27 static void block(struct sem *s)
28 {
29     int i=0;
30     sigset_t blk_mask, empty_mask;
31     sigaddset(&blk_mask, SIGUSR1);
32     sigemptyset(&empty_mask);
33
34     spin_lock(&s->countlock);
35     if(s->count>0)
36     {
37         spin_unlock(&s->countlock);
38         return;
39     }
40     spin_lock(&s->waitlock);
41     sigprocmask(SIG_BLOCK, &blk_mask, NULL);
```

```

42         if( pid_push(s, getpid())<0 )
43         {
44             fprintf(stderr,"pid_push(s,%d) error, stack too large\n",getpid());
45             exit(-1);
46         }
47         spin_unlock(&s->waitlock);
48         spin_unlock(&s->countlock);
49
50         sigsuspend(&empty_mask);
51     }
52
53     static void wake(struct sem *s)
54     {
55         pid_t pid;
56         spin_lock(&s->waitlock);
57         while( (pid=pid_pop(s))>0 )
58             kill(pid,SIGUSR1);
59         spin_unlock(&s->waitlock);
60     }
61
62     void sem_init(struct sem *s, int count)
63     {
64         s->count = count;
65         s->countlock = 0;
66         s->waitlock = 0;
67         s->waitlen = 0;
68         signal(SIGUSR1,handler);
69     }
70
71     int sem_try(struct sem *s)
72     {
73         spin_lock(&s->countlock);
74         if(s->count<1)
75         {
76             spin_unlock(&s->countlock);
77             return 0;
78         }
79         s->count--;
80         spin_unlock(&s->countlock);
81         return 1;
82     }
83
84     void sem_wait(struct sem *s)
85     {
86         while(!sem_try(s))
87             block(s);
88     }
89
90     void sem_inc(struct sem *s)
91     {
92         spin_lock(&s->countlock);
93         s->count++;
94         if(s->count==1)
95             wake(s);
96         spin_unlock(&s->countlock);
97     }

```

SPINLOCK

spinlib.h

```
1  #ifndef __SPINLIB_H
2  int spin_lock(volatile char *lock);
3  int spin_unlock(volatile char *lock);
4  #define __SPINLIB_H
5  #endif
```

spinlib.c

```
1  #include "spinlib.h"
2  #include "tas.h"
3  #include <sched.h>
4
5  int spin_lock(volatile char *lock)
6  {
7      while(tas(lock))
8          sched_yield();
9      return 1;
10 }
11 int spin_unlock(volatile char *lock)
12 {
13     *lock=0;
14     return 0;
15 }
```

TAS

tas.h

```
1  #ifndef __TAS_H
2  int tas(volatile char *lock);
3  #define __TAS_H
4  #endif
```

tas.s

```
1      .text
2  .global tas
3      .type    tas,@function
4  tas:
5      pushq    %rbp
6      movq     %rsp, %rbp
7      movq     $1, %rax
8  #APP
9      lock;xchgb %al,(<%rdi)
10 #NO_APP
11      movsbq   %al,%rax
12      pop     %rbp
13      ret
14 .Lfe1:
15      .size    tas, .Lfe1-tas
```

MAKEFILE

Makefile

```
1 spintest.out:
2     gcc -o spintest.out spintest.c spinlib.c tas.s
3 semtest.out:
4     gcc -o semtest.out semtest.c sem.c spinlib.c tas.s
5 fifotest.out:
6     gcc -o fifotest.out fifotest.c fifo.c sem.c spinlib.c tas.s
7 acidtest.out:
8     gcc -o acidtest.out acidtest.c fifo.c sem.c spinlib.c tas.s
```